

Name : Suryal D . Khirade

Roll NO: T190424399

Assignment No : 05

Data Analytics II

1. Implement logistic regression using Python/R to perform classification on

Social\_Network\_Ads.csv dataset.

2. Compute Confusion matrix to find TP, FP, TN, FN, Accuracy, Error rate, Precision, Recall on the given dataset.

## STEP 1

INCLUDE NECESSARY LIBRARIES

```
In [3]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

## STEP 2

upload dataset

```
In [5]: data=pd.read_csv("C:\\Users\\alisu\\Downloads\\Social_Network_Ads.csv")
```

```
In [6]: data
```

Out[6]:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
...	...	...	...	...	...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

400 rows × 5 columns

## STEP 3

### EXPLORATORY DATA ANALYSIS

In [8]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   User ID                400 non-null   int64  
1   Gender                 400 non-null   object  
2   Age                   400 non-null   int64  
3   EstimatedSalary        400 non-null   int64  
4   Purchased              400 non-null   int64  
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
```

In [9]: data.describe()

Out[9]:

	User ID	Age	EstimatedSalary	Purchased
count	4.000000e+02	400.000000	400.000000	400.000000
mean	1.569154e+07	37.655000	69742.500000	0.357500
std	7.165832e+04	10.482877	34096.960282	0.479864
min	1.556669e+07	18.000000	15000.000000	0.000000
25%	1.562676e+07	29.750000	43000.000000	0.000000
50%	1.569434e+07	37.000000	70000.000000	0.000000
75%	1.575036e+07	46.000000	88000.000000	1.000000
max	1.581524e+07	60.000000	150000.000000	1.000000

In [10]: data.isnull().sum()

Out[10]:

User ID	0
Gender	0
Age	0
EstimatedSalary	0
Purchased	0

dtype: int64

```
In [11]: data.duplicated()
```


```
Out[11]: 0      False
         1      False
         2      False
         3      False
         4      False
         ...
        395     False
        396     False
        397     False
        398     False
        399     False
        Length: 400, dtype: bool
```

## STEP 4

To find outliers

```
In [13]: plt.figure(figsize=(30,50))  
sns.boxplot(data)
```

```
Out[13]: <Axes: >
```



## STEP 5

Conversion of categorical variable to numerical variable

```
In [14]: data.Gender=data.Gender.replace({"Male":1,"Female":0})
```

```
In [15]: data
```

```
Out[15]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	1	19	19000	0
1	15810944	1	35	20000	0
2	15668575	0	26	43000	0
3	15603246	0	27	57000	0
4	15804002	1	19	76000	0
...	...	...	...	...	...
395	15691863	0	46	41000	1
396	15706071	1	51	23000	1
397	15654296	0	50	20000	1
398	15755018	1	36	33000	0
399	15594041	0	49	36000	1

400 rows × 5 columns

## Step 6

Splitting dependent and independent variables

```
In [16]: x=data.drop( "Purchased",axis="columns")
```

In [17]: x

Out[17]:

	User ID	Gender	Age	EstimatedSalary
0	15624510	1	19	19000
1	15810944	1	35	20000
2	15668575	0	26	43000
3	15603246	0	27	57000
4	15804002	1	19	76000
...	...	...	...	...
395	15691863	0	46	41000
396	15706071	1	51	23000
397	15654296	0	50	20000
398	15755018	1	36	33000
399	15594041	0	49	36000

400 rows × 4 columns

In [18]: y=data.Purchased

In [19]: y

Out[19]:

```
0      0
1      0
2      0
3      0
4      0
..
395    1
396    1
397    1
398    0
399    1
```

Name: Purchased, Length: 400, dtype: int64

## STEP 7

Split the dataset

In [20]: `from sklearn.model_selection import train_test_split
xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.2,random_state=3)`

```
In [21]: xtrain
```

Out[21]:

	User ID	Gender	Age	EstimatedSalary
239	15772073	0	53	143000
188	15674206	1	35	72000
240	15701537	1	42	149000
23	15599081	0	45	22000
343	15629739	0	47	51000
...	...	...	...	...
256	15609637	0	41	72000
131	15801247	1	33	31000
249	15753102	0	35	97000
152	15699247	1	31	76000
362	15768072	0	47	50000

320 rows × 4 columns

```
In [22]: ytrain
```

Out[22]:

239	1
188	0
240	1
23	1
343	1
...	..
256	0
131	0
249	1
152	0
362	1

Name: Purchased, Length: 320, dtype: int64



In [23]: xtest

Out[23]:

	User ID	Gender	Age	EstimatedSalary
<b>376</b>	15596984	0	46	74000
<b>16</b>	15733883	1	47	25000
<b>365</b>	15807525	0	59	29000
<b>82</b>	15709476	1	20	49000
<b>107</b>	15789863	1	27	89000
...	...	...	...	...
<b>246</b>	15638003	0	35	50000
<b>10</b>	15570769	0	26	80000
<b>115</b>	15689237	1	40	57000
<b>74</b>	15592877	1	32	18000
<b>194</b>	15689751	1	28	89000

80 rows × 4 columns

In [24]: ytest

Out[24]:

```
376    0
16     1
365    1
82     0
107    0
..
246    0
10     0
115    0
74     0
194    0
```

Name: Purchased, Length: 80, dtype: int64

## STEP 8

Use of logistic regression

In [25]: `from sklearn.linear_model import LogisticRegression`  
`model=LogisticRegression()`

In [26]: model

Out[26]:

```
▼ LogisticRegression
LogisticRegression()
```

```
In [27]: model.fit(xtrain,ytrain)
```

```
Out[27]: ▾ LogisticRegression  
LogisticRegression()
```

```
In [28]: y_predict=model.predict(xtest)
```

```
In [29]: y_predict
```

```
Out[29]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,  
                0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0,  
                0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0,  
                0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], dtype=int64)
```

```
In [30]: compare_result= pd.DataFrame(y_predict,ytest)
```

```
In [31]: compare_result.head(50)
```

Out[31]:

	0
Purchased	
	0 0
	1 0
	1 0
	0 0
	0 0
	0 0
	1 0
	0 0
	0 0
	0 0
	0 0
	0 0
	1 0
	0 0
	0 0
	0 0
	0 0
	0 0
	0 0
	1 1
	0 0
	0 0
	0 0
	0 0
	0 0
	0 1
	0 0
	0 0
	1 1
	0 0
	1 0
	1 1
	0 0
	1 1
	0 1
	0 0
	0 0
	1 1

	0
Purchased	
0	0
1	0
0	0
1	1
1	1
1	1
0	0
1	0
0	0
0	0
1	1
0	0
0	0

```
In [46]: from sklearn.metrics import confusion_matrix , classification_report , precision_score
from sklearn.metrics import accuracy_score , roc_curve , error_rate
```

```
In [38]: accuracy=accuracy_score(ytest,y_predict)
```

```
In [39]: accuracy
```

```
Out[39]: 0.8
```

```
In [40]: confusion_matrix(ytest,y_predict)
```

```
Out[40]: array([[52,  3],
               [13, 12]], dtype=int64)
```

```
In [42]: precision_score(ytest,y_predict)
```

```
Out[42]: 0.8
```


```
In [43]: f1_score(ytest,y_predict)
```

```
Out[43]: 0.6
```


```
In [44]: roc_auc_score(ytest,y_predict)
```

```
Out[44]: 0.7127272727272727
```

```
In [47]: fpr, tpr, _ = roc_curve(ytest, y_predict)
#create ROC curve
plt.plot(fpr,tpr)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



```
In [48]: auc= roc_auc_score(ytest,y_predict)
plt.plot(fpr, tpr, label="AUC="+str(auc))
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.legend(loc=4)
plt.show()
```



```
In [52]: cr= classification_report(ytest,y_predict)
print(cr)
```

	precision	recall	f1-score	support
0	0.80	0.95	0.87	55
1	0.80	0.48	0.60	25
accuracy			0.80	80
macro avg	0.80	0.71	0.73	80
weighted avg	0.80	0.80	0.78	80

```
In [55]: y_predict= model.predict(xtest)
print("Testing accuracy:", model.score(xtest,ytest)*100)
```

Testing accuracy: 80.0